

# CSE138 Lecture 5



this time:

- implementing causal delivery
- totally-ordered delivery

causal broadcast

- unicast/multicast/broadcast


if → - intro distributed snapshots



Both of the above executions violate causal message delivery.


Don't deliver a message until you've delivered all the messages that are in the causal history of that message.

Given messages  $m_1$  and  $m_2$  in an execution, if  $m_1$ 's send happens before  $m_2$ 's send and they have the same recipient,  $m_1$  is delivered before  $m_2$  on the recipient process.



## implementing causal message delivery

- idea: use VCs to track message sends only.



A protocol for implementing causal message delivery that works when all messages are broadcast messages.



(i.e., a causal broadcast protocol)

Ken Birman et al. 1991

The protocol:


- 1) If a message sent by a process  $P_i$  is delivered at a process  $P_j$ , increment  $P_j$ 's VC in the  $P_i$  position (the sender's position).
- 2) If a message is sent by a process increment the sender's position in the sender's VC, and include that VC along with the message.
- 3) A message sent by a process  $P_i$  only gets delivered at  $P_j$  if, for the message's vector clock  $VC_m$ :

- $VC_m[P_i] = VC_{receiver}[P_i] + 1$ , and
  - $VC_m[P_k] \leq VC_{receiver}[P_k]$
- for all  $k \neq i$
- VC on the message.



important note: we're only tracking message sends in these VCs.


more generally, any time you use any logical clock to track events, you must decide which events are important for your use case.



works fine (but is overkill)

for implementing FIFO delivery, too.

totally-ordered message delivery




[1,0,0,0]

[0,0,0,1]

neither is bigger; these are VCs of independent events

Totally-ordered message delivery:

Given messages  $m_1$  and  $m_2$  in an execution, if any process delivers  $m_1$  and then  $m_2$ , all processes delivering both deliver  $m_1$  first.



This execution violates FIFO (and therefore causal) message delivery but it respects totally-ordered message delivery!

message delivery guarantees

strongest ← → weakest


causal FIFO

total + causal


possible, but hard.

or  
point-to-point


- unicast messages -  
one sender, one receiver




- broadcast messages -  
one sender, everyone receives  
(including the sender!)



- multicast messages -  
one sender, multiple receivers  
(not necessarily everyone)





A global snapshot has to combine all processes' local snapshots.

since  $m$ 's receive event is the snapshot, its send event should be, too — but it isn't! (global)

What does it mean for a snapshot to be correct?

We want this property:

if an event  $e$  is in the snapshot, then all events that happen before  $e$  should be, too.

A snapshot satisfying this property is called a consistent global snapshot.

Classic algorithm for this:

→ Chandy-Lamport algorithm (1985)  
 mani Chandy