


CSE138 Lecture 6

this time:

- Chandy-Lamport snapshots

channel - a connection from one process to another





C_{12} - Channel from P_1 to P_2 - empty

C_{21} - Channel from P_2 to P_1 - $[m_2]$

Assume that messages in a channel are delivered reliably, and in FIFO order.

initiating a snapshot - one (or more!) processes does the following:

- Record its own state.
- Send a marker message out on all its outgoing channels.
- start recording the messages it receives on all its incoming channels.



Receiving a marker message

Two cases to be concerned with: either it's the first marker message we've ever seen, or ... it's not.


When process P_i receives a marker on C_{ki} :

If it's the first marker P_i has seen:

- P_i records its state.
- P_i marks channel C_{ki} as "empty". (i.e., don't pay attention to C_{ki} anymore).
- P_i sends marker messages out on all its outgoing channels C_{ij} .
- P_i starts recording incoming messages on all its incoming channels except C_{ki} .

channel from process P_k to process P_i


(TLDR: Record your state and start recording incoming messages.)




back to describing the algorithm:

- If P_i receives a marker message on channel C_{ki} and P_i has already seen a marker:
 - ↑ either sent or received!
- P_i stops recording on channel C_{ki} , and sets C_{ki} 's final state as the sequence of all the messages that arrived on C_{ki} since recording began.


We're done taking our snapshot. Does it "make sense"?




this is bad!



this is also bad!



this is OK!




this is also OK!


In general, the Chandy-Lamport algorithm only takes consistent snapshots.

A consistent snapshot is one where, if an event e is in the snapshot, all events that happened before e are in the snapshot too.

Lamport's happens-before relation (not physical time)



BTW, we have to send $O(N^2)$ marker messages when snapshotting a system with N processes.



What if P_2 had said, "I don't want to be the token holder anymore. P_1 , you take over." in message m .

In the C-L algorithm, channel recordings help with this situation. In particular, we record the fact that m was sent on C_{21} , so m isn't forgotten about!

Weird and wonderful thing about the C-L algorithm: it works fine with multiple initiators!

