


this time:

- ✓ - Chandy-Lamport wrap-up
 - ✓ - uses of snapshots
 - ✓ - centralized vs. decentralized algorithms
 - ✓ - cuts and Consistent Cuts
 - ✓ - safety and liveness
- if time {
- reliable delivery
 - fault models
-

channels - connect one process to another, FIFO behavior, reliable behavior




FIFO violation



won't happen
in channels

(sort of) reliable delivery violation



example of
a safety
property.

example of
a liveness
property.


what does it mean for a snapshot algorithm to be "correct"?

- safety {
- snapshots it takes are consistent snapshots.
 - (if an event e is in the snapshot, then all events e' such that $e' \rightarrow e$ are also in the snapshot.)
- liveness {
- termination! I want to eventually get done taking a snapshot.

Does the C-L algorithm terminate?

What assumptions do we have to make?

- Fixed number of processes



if every process does this stuff, we terminate!


a process has to record:

- its own process state ✓
- the states of its incoming channels. ✓


once everyone receives a marker message on all of their incoming channels, this is done.

in the C-L algorithm, at least one process has to initiate.

but more than one is OK!



This is actually a huge deal that multiple processes can initiate the algorithm independently. (because if this weren't true, we would have to solve a hard agreement problem to decide who gets to initiate).



if processes can crash, termination is not guaranteed. (e.g., here P₁ would be waiting forever to get a marker message that's never being sent.)

Because this algorithm works even with multiple initiators, it is an example of a decentralized algorithm.

A decentralized algorithm is one that can be independently initiated by different processes without them needing to coordinate with each other.

(By contrast, a centralized algorithm must have exactly one initiator.)

What are snapshots for?

- Checkpointing - take periodic snapshots of the state of a running application, so that in case of failure, they have a consistent state to start from

"as if nothing had gone wrong",

VLDB


2017 Flink paper

incl. deadlock detection

↑ according to the Apache Flink docs, anyway.


- Stable property detection -

to detect if a stable property is true of a system, take a snapshot and inspect it.




deadlock is an example of a stable property:


a property that, once true, remains true forever.



Cuts and consistent cuts



A cut is an imaginary line through an execution dividing its events into "past" and "future".



A "good" snapshot algorithm, such as the C-L algorithm, will only take snapshots that correspond to consistent cuts.


An inconsistent cut is one in which there exist events e and e' such that $e \rightarrow e'$, but e' is in the "past" set of events and e is in the "future" set of events.

A little more about reliable delivery.

a liveness property




Reliable delivery goes hand in hand with safety properties such as FIFO delivery.



How do we actually get reliable delivery?

idea: Keep sending until you get an ack.



reliable delivery protocol