


# CSE 138 Lecture 12

- ✓ - announcements
- ✓ - A3 stuff
- Paxos!



"Strong" (Linearizable)

Causal

FIFO

**RYW**

Stronger

Weaker

Let's talk Paxos!

family of  
Paxos - a consensus protocols  
invented by Leslie Lamport  
~~1990~~ 1998  
↑


"The Part-Time Parliament"

2001 - "Paxos Made Simple"

Deniz Altinbeken - "Paxos Made Moderately Complex"  
R. van Renesse 2015(?)

3 roles that a process can play:

- proposers propose values.
- acceptors contribute to choosing from the proposed values.
- learners learn the agreed-upon value.




A process can play multiple roles!

Any process that plays any role(s) is a Paxos node.

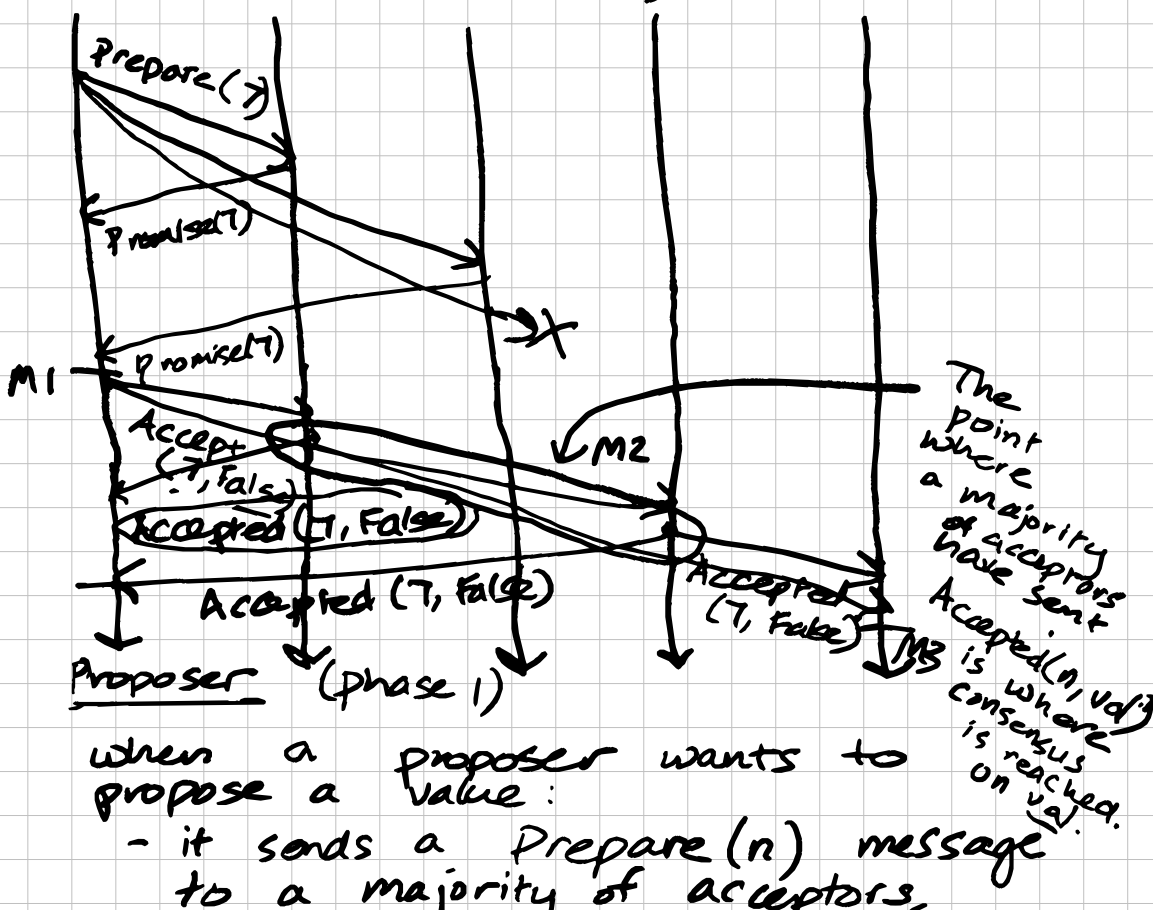
Requirements for a Paxos node:

- Paxos nodes have to know how many acceptors there are. In particular, they have to know what a majority of acceptors is.
- Paxos nodes have to be able to persist data. (in particular, they have to remember stuff like: what have I accepted!)

The version of Paxos we're discussing today is for deciding on a single value.



A run of Paxos! decide: True/False



when a proposer wants to propose a value:

- it sends a  $\text{Prepare}(n)$  message to a majority of acceptors, where  $n$  is:
  - higher than any proposal number this proposer has used before
  - unique to this proposer (no one else gets to use this proposal number)

Acceptor (phase 1)

When an acceptor receives a  $\text{prepare}(n)$  message, it looks at  $n$ , and asks:

"Did I previously promise to ignore requests with this proposal number?"

- if yes: ignores it!
- if no: it responds with  $\text{Promise}(n)$ , which means, "I will ignore any request with a proposal number lower than  $n$ ."

\* note: it's actually more subtle than this! (we'll come back to this)

Phase 1 ends, Phase 2 begins!

Proposer (phase 2) has received  $\text{Promise}(n)$  messages from a majority of acceptors for proposal number  $n$ .

To propose value  $\text{val}$ , it sends  $\text{Accept}(n, \text{val})$  to a majority of acceptors. \*\*

\*\* = actually more complicated than this

Acceptor (phase 2)

When an acceptor receives an  $\text{Accept}(n, \text{val})$  message, it asks:

"Did I previously promise to ignore this proposal number?"

- if yes: ignores it!
- if no: responds with  $\text{Accepted}(n, \text{val})$

and

also sends  $\text{Accepted}(n, \text{val})$  to all the Learners.

Short summary of phases

Phase 1

Proposer sends  $\text{Prepare}(n)$

Acceptors send  $\text{Promise}(n)$

Phase 2 (Proposer has heard from a majority of  $\text{Promise}(n)$ )

Proposer sends  $\text{Accept}(n, \text{val})$


Acceptors send  $\text{Accepted}(n, \text{val})$

By getting  $\text{Promise}(n)$  messages from a majority of acceptors for a given  $n$ , milestone 1: A proposer gets a majority of acceptors on board with it.

milestone 2 (Consensus is reached) when a majority of acceptors have sent  $\text{Accepted}(n, \text{val})$  for a particular  $\text{val}$ .

milestone 3 - happens individually on each learner. When the learner receives  $\text{Accepted}(n, \text{val})$  for a particular  $\text{val}$  from a majority of acceptors, that learner knows consensus has been reached.

Let's make things more interesting and have multiple proposers!



Phase 1 Acceptor behavior

When an acceptor gets  $\text{Prepare}(n)$ , it asks: "Did I previously promise to ignore this proposal number?"

- if yes: ignores it!
- if no: it asks: "Have I previously accepted anything?"

- if yes: responds to the proposer with  $\text{Promise}(n, (n_{\text{prev}}, \text{val}_{\text{prev}}))$

$n_{\text{prev}}$  = highest previously accepted proposal number

$\text{val}_{\text{prev}}$  = previously accepted value.

- if no:  $\text{Promise}(n)$ .